


Autoencoders

(unsupervised model)



A special type of feedforward neural network which does:

Encodes its input x_i into a hidden representation h

$$h = g(Wx_i + b) \quad \text{activation fn}$$

Decodes the input again from this hidden representation

$$\hat{x}_i = f(W^*h + c) \quad \text{output funct}^n$$

Why are we doing this?

We want to capture the most important characteristics of the input x_i .

The model is trained to minimize a loss function which ensures \hat{x}_i and x_i are close.


Case 1: $\dim(h) < \dim(x_i)$ (undercomplete auto-encoder)

If we are able to reconstruct \hat{x}_i perfectly from h then h is a loss free encoding of x_i .

- h captures all important characteristics of x_i .


Case 2: $\dim(h) > \dim(x_i)$ (overcomplete auto-encoder)

In this case, auto-encoder could learn a trivial encoding - by simply copying x_i .



Choices of f and g

Case 1: Inputs are binary
 $x_{ij} \in \{0, 1\}$



Since i/p is binary, we would want the o/p to also be binary.

What should f be?

- sigmoid: $\sigma(W^*\mathbf{h} + \mathbf{c})$

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

Case 2: Real inputs

$$x_{ij} \in \mathbb{R}$$

f : linear

usual choice of g : tanh, sigmoid.

Choice of loss function

Case 1: When input is real

(i) Cross entropy


~~(ii)~~ Mean Squared error

$$L(\theta) = \frac{1}{N} \sum_{i=1}^N \|x_i - \hat{x}_i\|_2^2 = \frac{1}{N} \sum_{i=1}^N \underbrace{\sum_{j=1}^d (x_{ij} - \hat{x}_{ij})^2}_{L_i(\theta)}$$

$$L(\theta) = \frac{1}{N} \sum_{i=1}^N L_i(\theta)$$

$$L_i(\theta) = \sum_{j=1}^d (x_{ij} - \hat{x}_{ij})^2$$

$$x_i \in \mathbb{R}^d \\ (x_1, x_2, \dots, x_d)$$



Case 2: When input is binary
use cross entropy.

$$L(\theta) = \frac{1}{N} \sum L_i(\theta)$$

$$L_i(\theta) = - \left(\sum_{j=1}^d x_{ij} \log \hat{x}_{ij} + (1 - x_{ij}) \log (1 - \hat{x}_{ij}) \right)$$

Train the NN using backpropagation with the objective

$$\min_{w, w^*, b, c} L(\theta)$$

Link to PCA

Encoder part of auto-encoder is equivalent to PCA if :

1. We use a linear encoder
2. We use a linear decoder
3. Use squared error loss
4. Normalize the inputs :

$$\tilde{x}_{ij} = \frac{1}{\sqrt{N}} (x_{ij} - \bar{x}_{ij})$$

Regularization

(i) l_2 - regularization

$$\min \mathcal{L}(\theta) + \lambda \|\theta\|_2^2$$

(ii) Weight Sharing


$$W^* = W^T$$

Other Types of Auto-encoders

1. De-noising Auto-encoders.

Corrupts the input data using some probabilistic process

$P(\tilde{x}_{ij} | x_{ij})$ before feeding it to the network



Consider binary input

$$\begin{array}{l|l} P(\tilde{x}_{ij}=0 | x_{ij}) = 1-p & P(\tilde{x}_{ij} = \text{not}(x_{ij}) | x_{ij}) = 1-p \\ P(\tilde{x}_{ij}=x_{ij} | x_{ij}) = p & P(\tilde{x}_{ij} = x_{ij} | x_{ij}) = p \end{array}$$

- It does not make sense for the network to memorize the input.
- Instead the model has capture good characteristics of the input.

2. Sparse Auto-encoder

A hidden neuron has values between 0 and 1.

activated when value is close to 1.

deactivated " " " " " 0.

A sparse auto-encoder tries to ensure that the neuron is deactivated most of the time.

Another way of doing regularization.