Let's revisit the task of autocompletion.

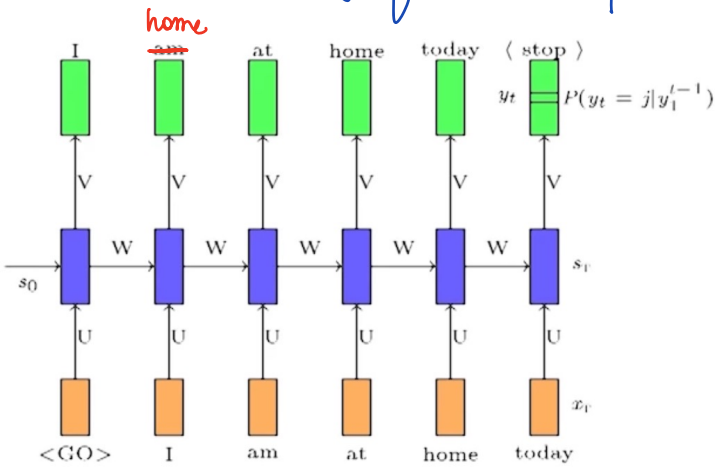

vocab

$= \{I, at, today, home, am\}$

$j=0$	0.1	I
$j=1$	0.1	at
$j=2$	0.2	today
$j=3$	0.1	home
$j=4$	0.5	am

At any time step t, we want to compute

$$\underset{j \in V}{\arg\max}\ P(y_t = j \mid y_{t-1}, y_{t-2}, \ldots, y_1)$$

j: word; V: Vocabulary

eg. $P(y_4 = \text{home} \mid \text{at}, \text{am}, \text{I})$

Using RNN,

$$P(y_t = j \mid y_{t-1}, y_{t-2}, \ldots, y_1) = (\text{softmax}(Vs_t + c))_j$$

s_t captures all the information until time step t.

At test time,

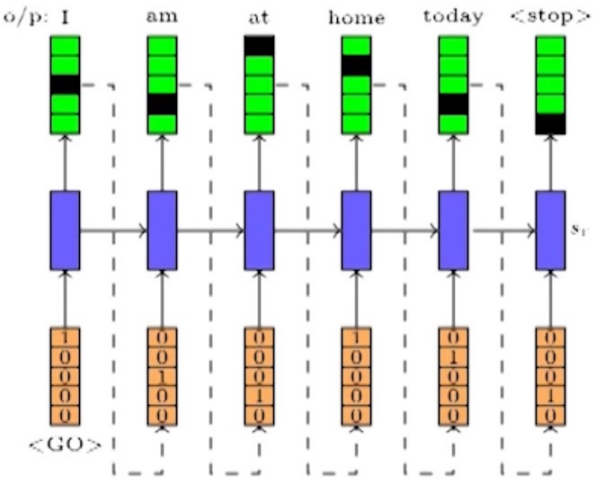

} one hot encoding of the words.

Image Captioning

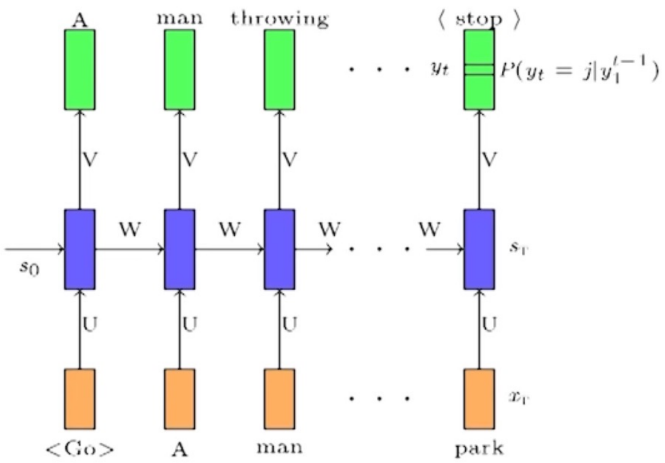

A man throwing (a) frisbee in a park

9:05 AM

We want to generate a sentence given an image. → Image.

We are interested in $P(y_t = j \mid y_{t-1}, \ldots, y_1, I)$

Since CNN architectures are good for images, we use it to learn important features from the image and then pass it on to the RNN model.

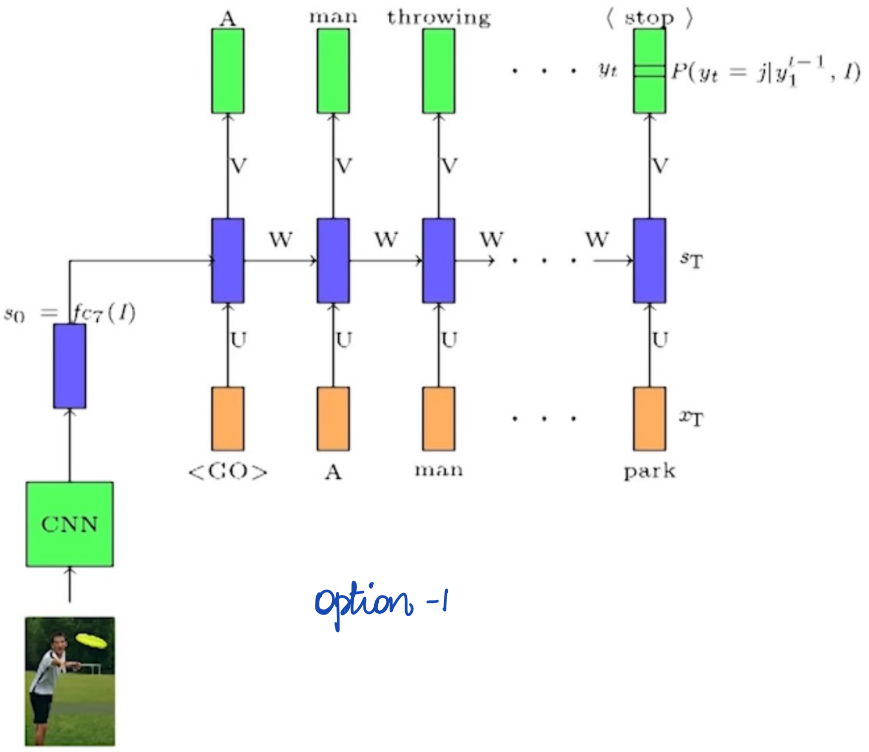

option -1

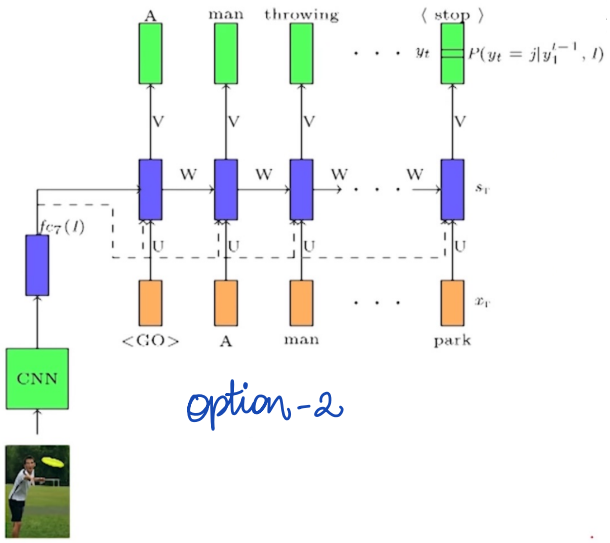

option-2

Encoder - Decoder

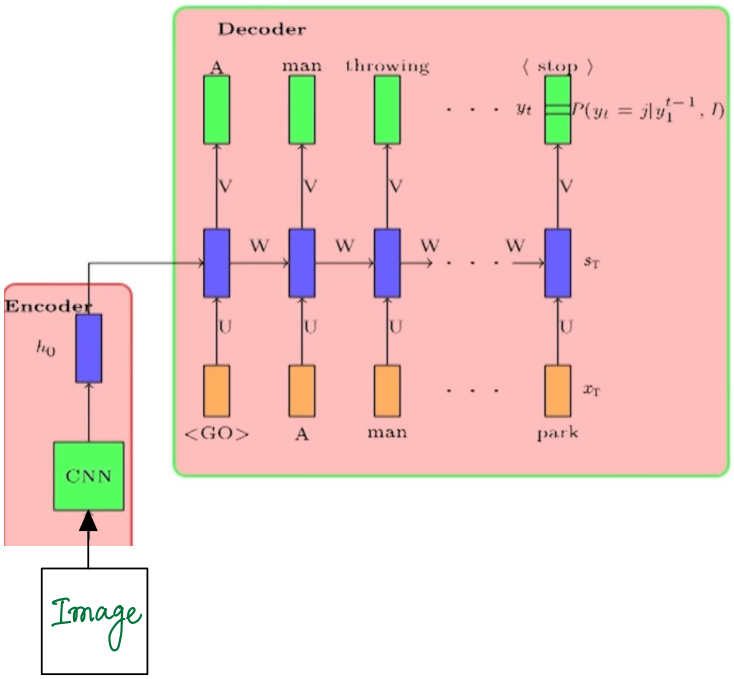

In image captioning,

A CNN is used to "encode" the image

- Learn good feature representation of the input.

An RNN is used to "decode" a sentence from this encoding.

Task: Image Captioning

Data: $\{ x_i = \text{image}, y_i = \text{caption} \}_{i=1}^{N}$

Model:

Encoder

$$s_0 = CNN(x_i)$$

Decoder

$$s_t = RNN(s_{t-1}, y_{t-1})$$

Parameters: U, V, W, b, c and all weights and biases of CNN.

Training all parameters of encoder and decoder together.

- End to end model

Machine Translation

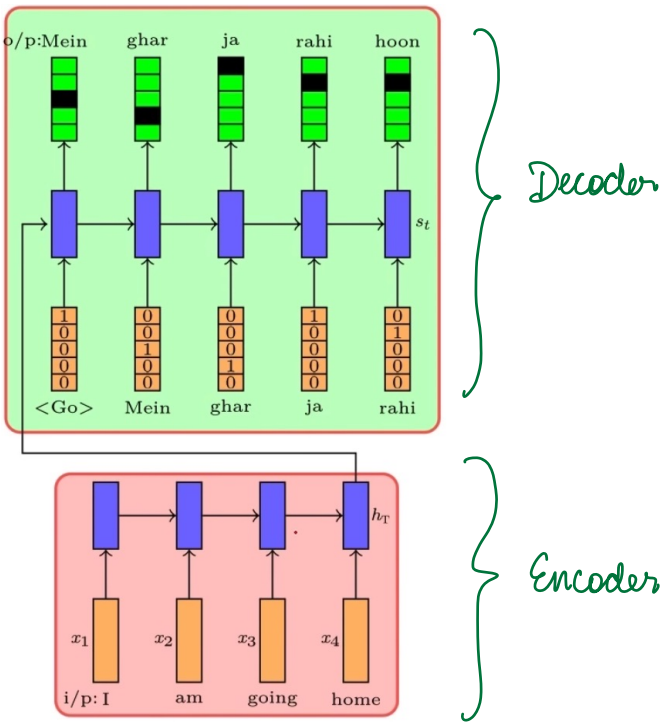

Video Captioning